

Annotation Guidelines

for our project on “Functional Tests for Hate Speech Detection Models”

Background:

We are looking to create a better way of evaluating hate speech detection models, by which we mean automated software for detecting hateful content. For this purpose, we have compiled a benchmark dataset consisting of a few thousand short text entries. Each entry reflects a particular aspect of online hate. We can use these entries to understand how hate speech detection models perform on the aspects we test.

Your tasks:

1. Mark up each entry assigned to you for whether it is **hateful** or **not hateful**.
2. Flag any entries which you think are **unrealistic**.

What do we mean by “hate”?

We define **hate** as *abuse that is targeted at a protected group or at its members for being a part of that group*. Protected groups are based on age, disability, gender reassignment, marriage and civil partnership, pregnancy and maternity, race (colour, nationality, ethnic or national origins), religion or belief, sex and sexual orientation as well as gender identity.

What do we mean by “unrealistic”?

We intentionally leave this up to you. We did not construct any entries to be unrealistic, so this optional flag is merely meant to ensure data quality.

Some things to keep in mind while annotating:

- Language that is hateful in some contexts can have non-hateful uses (e.g. homonyms of slurs: “*I smoked a fag with my mates*”, reclaimed slurs: “*I’m a bad bitch*”)
- The target of abuse matters for whether something is hateful (e.g. “*I hate Jews*” and “*I hate Asian people like you*” are hateful but “*I hate you*” and “*I hate doctors*” are not)

Other key bits of guidance:

1. Do not try to ‘game’ how much hate you think there is or should be (i.e. do not aim for a certain percentage of hate).
2. Use the realism flag only for cases you think are genuinely unrealistic, not just because you think they are likely to appear infrequently.
3. Please complete your annotations independently and don’t talk about them with others.
4. Most statements are quite short so please don’t overthink them.

If you have any technical issues or smaller questions, please DM us on Slack. In case of any more substantial concerns, please email us.

Thanks a lot for your work on this!